

GroupBy with Aggregate Functions

Single Column:

Use `groupBy("column_name")` and apply an aggregate function (e.g., `sum()`, `avg()`, `count()`).

```
df.groupBy("category").agg({"sales": "sum"}).show()
```

Multiple Columns:

Group by multiple columns by passing a list.

```
df.groupBy("category", "region").agg({"sales": "avg"}).show()
```

Pivot

Reshape data, turning unique values of a column into new column headers.

Used with aggregation.

```
df.groupBy("category").pivot("year").sum("sales").show()
```

Dropping Null Values

Use `dropna()` to remove rows containing null values.

Drop rows with nulls in **any column**:

```
df.dropna().show()
```

Drop rows with nulls in **specific columns**:

```
df.dropna(subset=["column1", "column2"]).show()
```

Drop rows based on a **threshold** of non-null values:

```
df.dropna(thresh=2).show() # At least 2 non-null values
```

Sorting in PySpark: sortBy and orderBy

sort() Sorts a **DataFrame** based on one or more columns.

Equivalent to orderBy in PySpark; they perform the same function.

```
df.sort("column1", "column2").show() # Default is ascending
```

```
df.sort(df["column1"].desc()).show() # Specify descending
```

orderBy() Another way to sort a DataFrame. Allows sorting by **multiple columns**.

Offers more flexibility by supporting column expressions for sorting.

```
df.orderBy("column1", "column2").show() # Default ascending
```

```
df.orderBy(df["column1"].desc(), df["column2"].asc()).show() # Mixed order
```

Types of Sorting

Single Column Sorting:

```
df.orderBy("column1").show()
```

```
df.orderBy(df["column1"].desc()).show()
```

Multiple Column Sorting:

Ascending for all: `df.orderBy("column1", "column2").show()`

Mixed order: `df.orderBy(df["column1"].asc(), df["column2"].desc()).show()`

Using Expressions: Sorting can involve transformations:

```
from pyspark.sql.functions import col
```

```
df.orderBy(col("column1") + 1).show() # Sort by column1 incremented by 1
```

DE138 SIVAPRAKASH V

create data

```

▶ ✓ 11:06 AM (1s)

from pyspark.sql import SparkSession

# Initialize SparkSession
spark = SparkSession.builder \
    .appName("example") \
    .getOrCreate()

# Data
simpleData = [("James", "Sales", "NY", 90000, 34, 10000),
              ("Michael", "Sales", "NY", 86000, 56, 20000),
              ("Robert", "Sales", "CA", 81000, 30, 23000),
              ("Maria", "Finance", "CA", 90000, 24, 23000),
              ("Raman", "Finance", "CA", 99000, 40, 24000),
              ("Scott", "Finance", "NY", 83000, 36, 19000),
              ("Jen", "Finance", "NY", 79000, 53, 15000),
              ("Jeff", "Marketing", "CA", 80000, 25, 18000),
              ("Kumar", "Marketing", "NY", 91000, 50, 21000)]

# Create DataFrame
schema = ["employee_name", "department", "state", "salary", "age", "bonus"]
df = spark.createDataFrame(data=simpleData, schema=schema)
df.show()

```

pivot()

```

▶ ✓ 11:06 AM (3s)

#using pivot function
df.groupBy("department").sum("salary").show()
df.groupBy("department").pivot("employee_name").sum("salary").show()

```

▶ (9) Spark Jobs

```

+-----+-----+
|department|sum(salary)|
+-----+-----+
|    Sales|    257000|
|   Finance|    351000|
|Marketing|    171000|
+-----+-----+

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
|department|James|Jeff|Jen|Kumar|Maria|Michael|Raman|Robert|Scott|
+-----+-----+-----+-----+-----+-----+-----+-----+
|    Sales|90000| null| null| null| null| 86000| null| 81000| null|
|   Finance| null| null|79000| null|90000| null|99000| null|83000|
|Marketing| null|80000| null|91000| null| null| null| null| null|
+-----+-----+-----+-----+-----+-----+-----+-----+

```

salary aggregate sum before and after grouping

```

11:06 AM (1s) 18

df.groupBy("department").agg({"salary": "sum"}).show()
df.agg({"salary": "sum"}).show() # Without group using agg on salary columns

(4) Spark Jobs

+-----+-----+
|department|sum(salary)|
+-----+-----+
|Sales|257000|
|Finance|351000|
|Marketing|171000|
+-----+-----+

+-----+
|sum(salary)|
+-----+
|779000|
+-----+

```

Working with Loan Data

Create data

```

11:51 AM (1s) 33

data_loan = spark.read.csv("/FileStore/tables/LoanData-2.csv", inferSchema=True, header=True)

(2) Spark Jobs

data_loan: pyspark.sql.dataframe.DataFrame = [Loan_ID: string, Gender: string ... 11 more fields]

```

data exploration

12:24 PM (-1s) 35 Python

```
data_loan.limit(10).display()
```

(1) Spark Jobs

	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
1	0	Graduate	No	5849	0	na	360
2	1	Graduate	No	4583	1508	128	360
3	0	Graduate	Yes	3000	0	66	360
4	0	Not Graduate	No	2583	2358	120	360
5	0	Graduate	No	6000	0	141	360
6	2	Graduate	Yes	5417	4196	267	360
7	0	Not Graduate	No	2333	1516	95	360
8	3+	Graduate	No	3036	2504	158	360
9	2	Graduate	No	4006	1526	168	360

```

11:59 AM (1s)

data_loan.printSchema()
print((data_loan.count(), len(data_loan.columns)))

▶ (2) Spark Jobs

root
|-- Loan_ID: string (nullable = true)
|-- Gender: string (nullable = true)
|-- Married: string (nullable = true)
|-- Dependents: string (nullable = true)
|-- Education: string (nullable = true)
|-- Self_Employed: string (nullable = true)
|-- ApplicantIncome: integer (nullable = true)
|-- CoapplicantIncome: double (nullable = true)
|-- LoanAmount: integer (nullable = true)
|-- Loan_Amount_Term: integer (nullable = true)
|-- Credit_History: integer (nullable = true)
|-- Property_Area: string (nullable = true)
|-- Loan_Status: string (nullable = true)

(614, 13)

```

groupby

```

12:05 PM (2s)

data_loan.groupBy("Education").min("ApplicantIncome").show()
data_loan.groupBy("Married").avg("CoapplicantIncome").show()
data_loan.groupBy("Loan_Status").count().show()

▶ (6) Spark Jobs

| Education|min(ApplicantIncome)|
+-----+-----+
|Not Graduate|          210|
| Graduate|          150|
+-----+-----+

+-----+-----+
|Married|avg(CoapplicantIncome)|
+-----+-----+
| null| 251.33333333333334|
| No| 1316.5586854460093|
| Yes| 1794.632964795578|
+-----+-----+

+-----+-----+
|Loan_Status|count|
+-----+-----+
| Y| 422|
| N| 192|
+-----+-----+

```

```

▶ 12:05 PM (2s) 39

data_loan.groupBy("Education", "Loan_Status", "Married").min("CoapplicantIncome").show()
data_loan.groupBy("Loan_Status", "Education").mean("ApplicantIncome").show()
data_loan.groupBy("Married", "Education").count().show()

▶ (6) Spark Jobs

+-----+-----+-----+
|Loan_Status| Education|avg(ApplicantIncome)|
+-----+-----+-----+
|          Y| Graduate| 5751.576470588236|
|          N|Not Graduate| 3646.4423076923076|
|          N| Graduate| 6114.5142857142855|
|          Y|Not Graduate| 3860.256097560976|
+-----+-----+-----+

+-----+-----+-----+
|Married| Education|count|
+-----+-----+-----+
|   Yes| Graduate| 309|
|   Yes|Not Graduate| 89|
|    No|Not Graduate| 45|
|    No| Graduate| 168|
|  null| Graduate| 3|
+-----+-----+-----+

```

drop null()

```

▶ 12:07 PM (1s) 41

loan_copy = data_loan
loan_clean = loan_copy.na.drop()

▶ (4) Spark Jobs

▶ loan_copy: pyspark.sql.dataframe.DataFrame = [Loan_ID: string, Gender: string ... 11 more fields]
▶ loan_clean: pyspark.sql.dataframe.DataFrame = [Loan_ID: string, Gender: string ... 11 more fields]

(614, 13)
(480, 13)

```

data after dropping null

```

+ Code + Text

▶ 12:08 PM (1s) 43

print((data_loan.count(), len(data_loan.columns)))
print((loan_clean.count(), len(loan_clean.columns)))

▶ (4) Spark Jobs

(614, 13)
(480, 13)

```

groupby after null drop

12:09 PM (4s)

45

```
loan_clean.groupBy("Education").min("ApplicantIncome").show()
loan_clean.groupBy("Married").avg("CoapplicantIncome").show()
loan_clean.groupBy("Loan_Status").count().show()

loan_clean.groupBy("Education", "Loan_Status", "Married").min("CoapplicantIncome").show()
loan_clean.groupBy("Loan_Status", "Education").mean("ApplicantIncome").show()
loan_clean.groupBy("Married", "Education").count().show()
```

(12) Spark Jobs

```
+-----+-----+
| Education|min(ApplicantIncome)|
+-----+-----+
|Not Graduate|      1442|
| Graduate|      150|
+-----+-----+

+-----+-----+
|Married|avg(CoapplicantIncome)|
+-----+-----+
| No| 1215.8934911242604|
| Yes| 1779.5463665229581|
+-----+-----+

+-----+-----+
|Loan_Status|count|
+-----+-----+
| Y| 332|
| N| 148|
+-----+-----+
```

(12) Spark Jobs

```
+-----+-----+-----+-----+
| Education|Loan_Status|Married|min(CoapplicantIncome)|
+-----+-----+-----+-----+
|Not Graduate| Y| No| 0.0|
| Graduate| N| Yes| 0.0|
|Not Graduate| N| Yes| 0.0|
|Not Graduate| N| No| 0.0|
| Graduate| N| No| 0.0|
|Not Graduate| Y| Yes| 0.0|
| Graduate| Y| Yes| 0.0|
| Graduate| Y| No| 0.0|
+-----+-----+-----+-----+

+-----+-----+-----+-----+
|Loan_Status| Education|avg(ApplicantIncome)|
+-----+-----+-----+-----+
| Y| Graduate| 5465.49446494465|
| N|Not Graduate| 3654.3333333333335|
| N| Graduate| 6397.428571428572|
| Y|Not Graduate| 4026.4590163934427|
+-----+-----+-----+-----+
```

```
+-----+-----+-----+
|Married| Education|count|
+-----+-----+-----+
| Yes| Graduate| 248|
| Yes|Not Graduate| 63|
| No|Not Graduate| 34|
| No| Graduate| 135|
+-----+-----+-----+
```

Joins

Data creation

```

from pyspark.sql import SparkSession

# Initialize SparkSession
spark = SparkSession.builder \
    .appName("example") \
    .getOrCreate()

# Data
emp = [(1,"Smith",-1,"2018","10","M",3000),(2, "Rose",1 , "2018", "20","M", 4000),(3,"Williams",1,"2010","10","M",1000),(4, "Jones",2 , "2005","10","F",2000),(5,"Brown",2,"2010","40","", -1),(6, "Brown", 2, "2018","50", "", -1)]
empColumns = ["emp_id","name","superior_emp_id","year_joined", "emp_dept_id","gender","salary"]

empDF = spark.createDataFrame(data=emp, schema = empColumns)
empDF.printSchema()
empDF.show()

dept = [("Finance",10),("Marketing",20),("Sales",30),("IT",40)]
deptColumns = ["dept_name","dept_id"]
deptDF = spark.createDataFrame(data=dept, schema = deptColumns)
deptDF.printSchema()
deptDF.show()

```

(6) Spark Jobs

```

empDF: pyspark.sql.dataframe.DataFrame = [emp_id long, name string, ... 5 more fields]
deptDF: pyspark.sql.dataframe.DataFrame = [dept_name string, dept_id long]

```

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
1	Smith	-1	2018	10	M	3000
2	Rose	1	2018	20	M	4000
3	Williams	1	2010	10	M	1000
4	Jones	2	2005	10	F	2000
5	Brown	2	2010	40		-1
6	Brown	2	2018	50		-1

```

root
 |-- dept_name: string (nullable = true)
 |-- dept_id: long (nullable = true)

```

dept_name	dept_id
Finance	10
Marketing	20
Sales	30
IT	40

Inner, Outer, Full

```

#Inner join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "inner").show()

```

(3) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2018	20	M	4000	Marketing	20
5	Brown	2	2010	40		-1	IT	40


```
#outer join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "outer").show()
```

▶ (3) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
null	null	null	null	null	null	null	Sales	30
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	null	null

```
#full join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "full").show()
```

▶ (3) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
null	null	null	null	null	null	null	Sales	30
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	null	null

left, left outer

▶ ✓ 12:40 PM (1s)

54

```
#Left join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "left").show()
```

▶ (6) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	null	null

```
#Left join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "leftouter").show()
```

▶ (6) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
5	Brown	2	2010	40		-1	IT	40
6	Brown	2	2010	50		-1	null	null

Right, Right Outer

▶ ✓ 12:41 PM (1s) 57

```
#right join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "right").show()
```

▶ (6) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
4	Jones	2	2005	10	F	2000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
null	null	null	null	null	null	null	Sales	30
5	Brown	2	2010	40		-1	IT	40

```
#right outer join
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id, "rightouter").show()
```

▶ (6) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary	dept_name	dept_id
4	Jones	2	2005	10	F	2000	Finance	10
3	Williams	1	2010	10	M	1000	Finance	10
1	Smith	-1	2018	10	M	3000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
null	null	null	null	null	null	null	Sales	30
5	Brown	2	2010	40		-1	IT	40

LeftSemi LeftAnti

▶ ✓ 12:42 PM (1s)

```
#leftsemi join
empDF.join(deptDF, empDF.emp_dept_id == deptDF.dept_id, "leftsemi").show()
```

▶ (3) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
1	Smith	-1	2018	10	M	3000
3	Williams	1	2010	10	M	1000
4	Jones	2	2005	10	F	2000
2	Rose	1	2010	20	M	4000
5	Brown	2	2010	40		-1

▶ ✓ 12:42 PM (2s)

```
#leftanti
empDF.join(deptDF, empDF.emp_dept_id == deptDF.dept_id, "leftanti").show()
```

▶ (6) Spark Jobs

emp_id	name	superior_emp_id	year_joined	emp_dept_id	gender	salary
6	Brown	2	2010	50		-1

Spark SQL using Loan Data (Clean)

Create View

```
06:45 PM (<1s) 64
loan_clean.createOrReplaceTempView("loan")
loan_clean.printSchema()
```

```
root
|-- Loan_ID: string (nullable = true)
|-- Gender: string (nullable = true)
|-- Married: string (nullable = true)
|-- Dependents: string (nullable = true)
|-- Education: string (nullable = true)
|-- Self_Employed: string (nullable = true)
|-- ApplicantIncome: integer (nullable = true)
|-- CoapplicantIncome: double (nullable = true)
|-- LoanAmount: integer (nullable = true)
|-- Loan_Amount_Term: integer (nullable = true)
|-- Credit_History: integer (nullable = true)
|-- Property_Area: string (nullable = true)
|-- Loan_Status: string (nullable = true)
```

Select query

06:47 PM (1s) 66 Python

```
spark.sql("select * from loan").limit(10).display()
loan_clean.select("Loan_ID", "ApplicantIncome", "Loan_Status").show(5)
```

(2) Spark Jobs

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	
5	LP001011	Male	Yes	2	Graduate	Yes	5417	
6	LP001013	Male	Yes	0	Not Graduate	No	2333	
7	LP001014	Male	Yes	3+	Graduate	No	3036	
8	LP001018	Male	Yes	2	Graduate	No	4006	
9	LP001020	Male	Yes	1	Graduate	No	12841	
10	LP001024	Male	Yes	2	Graduate	No	3200	

```
+-----+-----+
| Loan_ID|ApplicantIncome|Loan_Status|
+-----+-----+
| LP001003|      4583|      N|
| LP001005|      3000|      Y|
| LP001006|      2583|      Y|
| LP001008|      6000|      Y|
| LP001011|      5417|      Y|
+-----+-----+
only showing top 5 rows
```

filter based on data

06:49 PM (<1s) 68 Python

```
spark.sql("""SELECT * From loan WHERE ApplicantIncome<5000""").limit(10).display()
```

(1) Spark Jobs

id	Loan_ID	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
2	3000	0	66	360	1	Urban	Y
3	2583	2358	120	360	1	Urban	Y
4	2333	1516	95	360	1	Urban	Y
5	3036	2504	158	360	0	Semiurban	N
6	4006	1526	168	360	1	Urban	Y
7	3200	700	70	360	1	Urban	Y
8	3073	8106	200	360	1	Urban	Y
9	1853	2840	114	360	1	Rural	N
10	1299	1086	17	120	1	Urban	Y

10 rows | 0.33 seconds runtime Refreshed 8 minutes ago

order by

06:52 PM (<1s) 70 Python

```
spark.sql("""select * FROM loan WHERE Dependents in ('3+', '2')order by Dependents """).limit(10).display()
```

(1) Spark Jobs

id	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
2	LP001491	Male	Yes	2	Graduate	Yes	3316	
3	LP001245	Male	Yes	2	Not Graduate	Yes	1875	
4	LP001018	Male	Yes	2	Graduate	No	4006	
5	LP001267	Female	Yes	2	Graduate	No	1378	
6	LP001024	Male	Yes	2	Graduate	No	3200	
7	LP001318	Male	Yes	2	Graduate	No	6250	
8	LP001030	Male	Yes	2	Graduate	No	1299	
9	LP001319	Male	Yes	2	Not Graduate	No	3273	
10	LP001179	Male	Yes	2	Graduate	No	4616	

10 rows | 0.42 seconds runtime Refreshed 5 minutes ago

group by

4 minutes ago (1s) 72 Python

```
spark.sql("""SELECT Loan_Amount_Term,count(*) as count FROM loan GROUP BY Loan_Amount_Term""").show()
```

(2) Spark Jobs

Loan_Amount_Term	count
300	9
360	411
120	3
84	3
240	2
480	12
60	2
180	36
36	2