

Coding Challenge

1. Extract, Transform, Load in Pyspark.

ETL (Extract, Transform, Load) is a process used to prepare and integrate data for analysis and storage. In PySpark, ETL is performed on large datasets by using Spark's parallel processing capabilities. Here we can make use of clusters to perform the activity.

Extract involves retrieving data from various sources such as databases, APIs, or flat files. PySpark's built-in connectors and libraries, such as `spark.read`, support various file formats like CSV, JSON and Parquet for seamless extraction.

Transform is the core step where raw data is cleaned, enriched, and reshaped to meet business requirements. PySpark's DataFrame API and SQL capabilities enable transformations such as filtering, aggregations, joins, and column operations. The distributed architecture ensures transformations are fast and scalable, even for massive datasets.

Load involves storing the processed data into a target system, such as a database or a data warehouse, using PySpark's write functions. This completes the data pipeline, ensuring prepared data is readily available for analysis.

2. Queries & Solutions on Credit Dataset.

Read me:

For every query first I have written the Spark SQL code followed by PySpark code. question will be written in markdown.

1. First solution is on Spark SQL. For performing this initially I have created temporary view table. Then only we can perform SQL operations.
2. Second solution is on pyspark. I have performed these operations on the created dataframe itself.

For joins I have taken the credit data and keeping CustomerID column as common column I have splitted the dataset into 2 dataframes namely df1 and df2. The schema for the same can be viewed in page number 5 [under 4. JOINS division]. Similar to above to perform SQL operations in it I have created views for the 2 dataframes namely table1 and table2. Then performed joins operations in it.

Coding Challenge -Sivaprakash V

Import libraries & Initiate session

```

▶ ✓ 04:10 PM (<1s) 3

# initialize the session
from pyspark import SparkContext
from pyspark.sql import SparkSession

sc = SparkContext.getOrCreate()
spark = SparkSession.builder.appName('Coding challenge').getOrCreate()

```

Loading Data

```

▶ ✓ 04:53 PM (2s) 5

data_credit = spark.read.csv("/FileStore/tables/creditCard.csv", inferSchema=True, header=True)
df2 = spark.read.csv("/FileStore/tables/bankcredit.csv", inferSchema=True, header=True)
df1 = spark.read.csv("/FileStore/tables/info.csv", inferSchema=True, header=True)

```

▶ (6) Spark Jobs

```

▶ data_credit: pyspark.sql.dataframe.DataFrame = [RowNumber: integer, CustomerId: integer ... 11 more fields]
▶ df2: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Tenure: integer ... 5 more fields]
▶ df1: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 4 more fields]

```

Data Exploration

▶ ✓ 04:10 PM (<1s)

```

# Print Schema
data_credit.printSchema()

```

```

root
|-- RowNumber: integer (nullable = true)
|-- CustomerId: integer (nullable = true)
|-- Surname: string (nullable = true)
|-- CreditScore: integer (nullable = true)
|-- Geography: string (nullable = true)
|-- Gender: string (nullable = true)
|-- Age: integer (nullable = true)
|-- Tenure: integer (nullable = true)
|-- Balance: double (nullable = true)
|-- NumOfProducts: integer (nullable = true)
|-- IsActiveMember: integer (nullable = true)
|-- EstimatedSalary: double (nullable = true)
|-- Exited: integer (nullable = true)

```

Question & Solution

Create View

```

04:21 PM (<1s)

# Creating Views
data_credit.createOrReplaceTempView('credit')

```

1. Filter out customers with a CreditScore greater than 700 and who are active members:

```

04:24 PM (1s) 16

spark.sql("SELECT * FROM credit WHERE CreditScore > 700 AND IsActiveMember = 1;").show()

```

(1) Spark Jobs

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	79084.1	0
7	15592531	Bartlett	822	France	Male	50	7	0.0	2	1	10062.8	0
20	15568982	Hao	726	France	Female	24	6	0.0	2	1	54724.03	0
21	15577657	McDonald	732	France	Male	41	8	0.0	2	1	170886.17	0
25	15625047	Yen	846	France	Female	38	5	0.0	1	1	187616.16	0
27	15736816	Young	756	Germany	Male	36	2	136815.64	1	1	170041.95	0
35	15732963	Clements	722	Spain	Female	29	9	0.0	2	1	142033.07	0
38	15729599	Lorenzo	804	Spain	Male	33	7	76548.6	1	1	98453.45	0
39	15717426	Armstrong	850	France	Male	36	7	0.0	1	1	40812.9	0
46	15754849	Tyler	776	Germany	Female	32	4	109421.13	2	1	126517.46	0
47	15602280	Martin	829	Germany	Female	27	9	112045.67	1	1	119708.21	1
66	15789484	Hammond	751	Germany	Female	36	6	169831.46	2	1	27758.36	0
68	15641582	Chibugo	735	Germany	Male	43	10	123180.01	2	1	196673.28	0
79	15575185	Bushell	757	Spain	Male	33	5	77253.22	1	1	194239.63	0
94	15640635	Capon	769	France	Male	29	8	0.0	2	1	172290.61	0
95	15676966	Capon	730	Spain	Male	42	4	0.0	2	1	85982.47	0
97	15738721	Graham	773	Spain	Male	41	9	102827.44	1	1	64595.25	0
98	15693683	Yuille	814	Germany	Male	29	8	97086.4	2	1	197276.13	0

```

04:26 PM (<1s) 17

filtered_df = data_credit.filter((data_credit.CreditScore > 700) & (data_credit.IsActiveMember == 1))
filtered_df.show()

```

(1) Spark Jobs

```

filtered_df: pyspark.sql.dataframe.DataFrame = [RowNumber: integer, CustomerId: integer ... 11 more fields]

```

20	15568982	Hao	726	France	Female	24	6	0.0	2	1	54724.03	0
21	15577657	McDonald	732	France	Male	41	8	0.0	2	1	170886.17	0
25	15625047	Yen	846	France	Female	38	5	0.0	1	1	187616.16	0
27	15736816	Young	756	Germany	Male	36	2	136815.64	1	1	170041.95	0
35	15732963	Clements	722	Spain	Female	29	9	0.0	2	1	142033.07	0
38	15729599	Lorenzo	804	Spain	Male	33	7	76548.6	1	1	98453.45	0
39	15717426	Armstrong	850	France	Male	36	7	0.0	1	1	40812.9	0
46	15754849	Tyler	776	Germany	Female	32	4	109421.13	2	1	126517.46	0
47	15602280	Martin	829	Germany	Female	27	9	112045.67	1	1	119708.21	1
66	15789484	Hammond	751	Germany	Female	36	6	169831.46	2	1	27758.36	0
68	15641582	Chibugo	735	Germany	Male	43	10	123180.01	2	1	196673.28	0
79	15575185	Bushell	757	Spain	Male	33	5	77253.22	1	1	194239.63	0
94	15640635	Capon	769	France	Male	29	8	0.0	2	1	172290.61	0
95	15676966	Capon	730	Spain	Male	42	4	0.0	2	1	85982.47	0
97	15738721	Graham	773	Spain	Male	41	9	102827.44	1	1	64595.25	0
98	15693683	Yuille	814	Germany	Male	29	8	97086.4	2	1	197276.13	0
108	15812878	Parsons	785	Germany	Female	36	2	99806.85	1	1	36976.52	0
115	15609618	Fanucci	721	Germany	Male	28	9	154475.54	2	1	101300.94	1

only showing top 20 rows

2. Calculate the average Balance, total EstimatedSalary, and count of customers:

04:28 PM (1s) 19

```
spark.sql("SELECT AVG(Balance) AS AverageBalance, SUM(EstimatedSalary) AS TotalSalary, COUNT(*) AS CustomerCount FROM credit;").display()
```

(2) Spark Jobs

Table +

	1.2 AverageBalance	1.2 TotalSalary	1.2 CustomerCount
1	76485.88928799961	1000902398.8099979	10000

04:30 PM (1s) 20

```
aggregations = data_credit.agg(
    {"Balance": "avg", "EstimatedSalary": "sum", "": "count"}
)
aggregations.withColumnRenamed("avg(Balance)", "AverageBalance").withColumnRenamed("sum(EstimatedSalary)", "TotalSalary") \
    .withColumnRenamed("count(1)", "CustomerCount").display()
```

(2) Spark Jobs

aggregations: pyspark.sql.dataframe.DataFrame = [count(1): long, avg(Balance): double ... 1 more field]

Table +

	1.2 CustomerCount	1.2 AverageBalance	1.2 TotalSalary
1	10000	76485.88928799961	1000902398.8099979

3. Find the average CreditScore and total number of customers grouped by Geography and Gender:

04:32 PM (1s) 22

```
spark.sql("SELECT Geography, Gender, AVG(CreditScore) AS AvgCreditScore, COUNT(*) AS CustomerCount FROM credit GROUP BY Geography, Gender;").display()
```

(2) Spark Jobs

Table +

	1.2 Geography	1.2 Gender	1.2 AvgCreditScore	1.2 CustomerCount
1	Germany	Female	653.0938809723386	1193
2	France	Male	650.0646567381038	2753
3	France	Female	649.1857585139319	2261
4	Spain	Male	650.9920749279539	1388
5	Germany	Male	649.9665653495441	1316
6	Spain	Female	651.7695133149679	1089

6 rows | 0.91 seconds runtime

04:33 PM (1s) 23

```
grouped_df = data_credit.groupBy("Geography", "Gender").agg({"CreditScore": "avg", "": "count"})
grouped_df.withColumnRenamed("avg(CreditScore)", "AvgCreditScore").withColumnRenamed("count(1)", "CustomerCount").display()
```

(2) Spark Jobs

grouped_df: pyspark.sql.dataframe.DataFrame = [Geography: string, Gender: string ... 2 more fields]

Table +

	1.2 Geography	1.2 Gender	1.2 CustomerCount	1.2 AvgCreditScore
1	Germany	Female	1193	653.0938809723386
2	France	Male	2753	650.0646567381038
3	France	Female	2261	649.1857585139319
4	Spain	Male	1388	650.9920749279539
5	Germany	Male	1316	649.9665653495441
6	Spain	Female	1089	651.7695133149679

6 rows | 0.84 seconds runtime

4. Joins

▶ ✓ 04:53 PM (<1s)

```
df1.printSchema()  
df2.printSchema()
```

```
root  
|-- CustomerId: integer (nullable = true)  
|-- Surname: string (nullable = true)  
|-- CreditScore: integer (nullable = true)  
|-- Geography: string (nullable = true)  
|-- Gender: string (nullable = true)  
|-- Age: integer (nullable = true)  
  
root  
|-- CustomerId: integer (nullable = true)  
|-- Tenure: integer (nullable = true)  
|-- Balance: double (nullable = true)  
|-- NumOfProducts: integer (nullable = true)  
|-- IsActiveMember: integer (nullable = true)  
|-- EstimatedSalary: double (nullable = true)  
|-- Exited: integer (nullable = true)
```

▶ ✓ 04:54 PM (<1s)

```
df1.createOrReplaceTempView("table1")  
df2.createOrReplaceTempView("table2")
```


4.1. Inner Join

04:57 PM (1s)

28

```
result = spark.sql("SELECT * FROM table1 t1 INNER JOIN table2 t2 ON t1.CustomerId = t2.CustomerId;")
result.show()
```

(2) Spark Jobs

```
result: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 11 more fields]
```

CustomerId	Surname	CreditScore	Geography	Gender	Age	CustomerId	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
15634602	Hargrave	619	France	Female	42	15634602	2	null	1	1	101348.88	1
15647311	Hill	608	null	Female	41	15647311	1	83807.86	1	1	112542.58	0
15701354	null	null	France	null	null	15701354	1	null	2	0	93826.63	0
15737888	Mitchell	850	Spain	Female	null	15737888	2	125510.82	1	1	79084.1	null
15574012	null	645	null	Male	44	15574012	8	113755.78	2	null	149756.71	1
15592531	Bartlett	822	France	Male	null	15592531	7	null	2	1	10062.8	0
15656148	Obinna	376	Germany	null	29	15656148	4	115046.74	4	0	119346.88	1
15792365	He	501	France	null	44	15792365	4	142051.07	2	1	74940.5	0
15592389	null	684	France	Male	27	15592389	2	134603.88	1	1	71725.73	null
15737452	null	null	Germany	Male	58	15737452	null	132602.88	1	null	5097.67	1
15788218	Henderson	549	Spain	Female	null	15788218	9	null	2	1	14406.41	0
15661507	Muldrow	587	Spain	Male	45	15661507	6	null	1	0	158684.81	0
15568982	Hao	726	France	Female	24	15568982	6	null	2	1	54724.03	0

04:54 PM (<1s)

29

```
inner_join = df1.join(df2, on="CustomerId", how="inner")
inner_join.show()
```

(2) Spark Jobs

```
inner_join: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 10 more fields]
```

CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
15634602	Hargrave	619	France	Female	42	2	null	1	1	101348.88	1
15647311	Hill	608	null	Female	41	1	83807.86	1	1	112542.58	0
15701354	null	null	France	null	null	1	null	2	0	93826.63	0
15737888	Mitchell	850	Spain	Female	null	2	125510.82	1	1	79084.1	null
15574012	null	645	null	Male	44	8	113755.78	2	null	149756.71	1
15592531	Bartlett	822	France	Male	null	7	null	2	1	10062.8	0
15656148	Obinna	376	Germany	null	29	4	115046.74	4	0	119346.88	1
15792365	He	501	France	null	44	4	142051.07	2	1	74940.5	0
15592389	null	684	France	Male	27	2	134603.88	1	1	71725.73	null
15737452	null	null	Germany	Male	58	null	132602.88	1	null	5097.67	1
15788218	Henderson	549	Spain	Female	null	9	null	2	1	14406.41	0
15661507	Muldrow	587	Spain	Male	45	6	null	1	0	158684.81	0
15568982	Hao	726	France	Female	24	6	null	2	1	54724.03	0

4.2. Left join

04:57 PM (1s)

31

```
result1 = spark.sql("SELECT * FROM table1 t1 LEFT JOIN table2 t2 ON t1.CustomerId = t2.CustomerId;")
result1.show()
```

(2) Spark Jobs

```
result1: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 11 more fields]
```

CustomerId	Surname	CreditScore	Geography	Gender	Age	CustomerId	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
15634602	Hargrave	619	France	Female	42	15634602	2	null	1	1	101348.88	1
15647311	Hill	608	null	Female	41	15647311	1	83807.86	1	1	112542.58	0
15701354	null	null	France	null	null	15701354	1	null	2	0	93826.63	0
15737888	Mitchell	850	Spain	Female	null	15737888	2	125510.82	1	1	79084.1	null
15574012	null	645	null	Male	44	15574012	8	113755.78	2	null	149756.71	1
15592531	Bartlett	822	France	Male	null	15592531	7	null	2	1	10062.8	0
15656148	Obinna	376	Germany	null	29	15656148	4	115046.74	4	0	119346.88	1
15792365	He	501	France	null	44	15792365	4	142051.07	2	1	74940.5	0
15592389	null	684	France	Male	27	15592389	2	134603.88	1	1	71725.73	null
15737452	null	null	Germany	Male	58	15737452	null	132602.88	1	null	5097.67	1
15788218	Henderson	549	Spain	Female	null	15788218	9	null	2	1	14406.41	0
15661507	Muldrow	587	Spain	Male	45	15661507	6	null	1	0	158684.81	0
15568982	Hao	726	France	Female	24	15568982	6	null	2	1	54724.03	0

04:54 PM (1s)

32

```
left_join = df1.join(df2, on="CustomerId", how="left")
left_join.show()
```

(2) Spark Jobs

```
left_join: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 10 more fields]
```

CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
15634602	Hargrave	619	France	Female	42	2	null	1	1	101348.88	1
15647311	Hill	608	null	Female	41	1	83807.86	1	1	112542.58	0
15701354	null	null	France	null	null	1	null	2	0	93826.63	0
15737888	Mitchell	850	Spain	Female	null	2	125510.82	1	1	79084.1	null
15574012	null	645	null	Male	44	8	113755.78	2	null	149756.71	1
15592531	Bartlett	822	France	Male	null	7	null	2	1	10062.8	0
15656148	Obinna	376	Germany	null	29	4	115046.74	4	0	119346.88	1
15792365	He	501	France	null	44	4	142051.07	2	1	74940.5	0
15592389	null	684	France	Male	27	2	134603.88	1	1	71725.73	null
15737452	null	null	Germany	Male	58	null	132602.88	1	null	5097.67	1
15788218	Henderson	549	Spain	Female	null	9	null	2	1	14406.41	0
15661507	Muldrow	587	Spain	Male	45	6	null	1	0	158684.81	0
15568982	Hao	726	France	Female	24	6	null	2	1	54724.03	0

4.3. Right Join

```

04:58 PM (1s) 34

result2 = spark.sql("SELECT * FROM table1 t1 RIGHT JOIN table2 t2 ON t1.CustomerId = t2.CustomerId;")
result2.show()

(2) Spark Jobs

result2: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 11 more fields]
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 15647311| Hill| 608| null|Female| 41| 15647311| 1| 83807.86| 1| 1| 112542.58| 0|
| null| null| null| null| null| null| 15619304| 8| 159660.8| 3| 0| 113931.57| 1|
| 15701354| null| null| France| null| null| 15701354| 1| null| 2| 0| 93826.63| 0|
| 15737888| Mitchell| 850| Spain|Female| null| 15737888| 2| 125510.82| 1| 1| 79084.1| null|
| 15574012| null| 645| null| Male| 44| 15574012| 8| 113755.78| 2| null| 149756.71| 1|
| 15592531| Bartlett| 822| France| Male| null| 15592531| 7| null| 2| 1| 10062.8| 0|
| 15656148| Obinna| 376| Germany| null| 29| 15656148| 4| 115046.74| 4| 0| 119346.88| 1|
| 15792365| He| 501| France| null| 44| 15792365| 4| 142051.07| 2| 1| 74940.5| 0|
| 15592389| null| 684| France| Male| 27| 15592389| 2| 134603.88| 1| 1| 71725.73| null|
| null| null| null| null| null| null| 15767821| null| 102016.72| 2| 0| 80181.12| 0|
| null| null| null| null| null| null| 15737173| 3| null| 2| 0| 76390.01| 0|
| null| null| null| null| null| null| 15632264| 10| null| 2| null| 26260.98| 0|
| null| null| null| null| null| null| 15691483| 5| null| 2| 0| 190857.79| null|
| null| null| null| null| null| null| 15600882| 7| null| 2| 1| 65951.65| 0|
| null| null| null| null| null| null| 15643966| 3| 143129.41| 2| 1| 64327.26| null|
| 15737452| null| null| Germany| Male| 58| 15737452| null| 132602.88| 1| null| 5097.67| 1|
| 15788218| Henderson| 549| Spain|Female| null| 15788218| 9| null| 2| 1| 14406.41| 0|
| 15661507| Muldrow| 587| Spain| Male| 45| 15661507| 6| null| 1| 0| 158684.81| 0|
| 15568982| Hao| 726| France|Female| 24| 15568982| 6| null| 2| 1| 54724.03| 0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

```

04:55 PM (1s) 35

right_join = df1.join(df2, on="CustomerId", how="right")
right_join.show()

(2) Spark Jobs

right_join: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 10 more fields]
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 15647311| Hill| 608| null|Female| 41| 1| 83807.86| 1| 1| 112542.58| 0|
| 15619304| null| null| null| null| null| 8| 159660.8| 3| 0| 113931.57| 1|
| 15701354| null| null| France| null| null| 1| null| 2| 0| 93826.63| 0|
| 15737888| Mitchell| 850| Spain|Female| null| 2| 125510.82| 1| 1| 79084.1| null|
| 15574012| null| 645| null| Male| 44| 8| 113755.78| 2| null| 149756.71| 1|
| 15592531| Bartlett| 822| France| Male| null| 7| null| 2| 1| 10062.8| 0|
| 15656148| Obinna| 376| Germany| null| 29| 4| 115046.74| 4| 0| 119346.88| 1|
| 15792365| He| 501| France| null| 44| 4| 142051.07| 2| 1| 74940.5| 0|
| 15592389| null| 684| France| Male| 27| 2| 134603.88| 1| 1| 71725.73| null|
| 15767821| null| null| null| null| null| null| 102016.72| 2| 0| 80181.12| 0|
| 15737173| null| null| null| null| null| 3| null| 2| 0| 76390.01| 0|
| 15632264| null| null| null| null| null| 10| null| 2| null| 26260.98| 0|
| 15691483| null| null| null| null| null| 5| null| 2| 0| 190857.79| null|
| 15600882| null| null| null| null| null| 7| null| 2| 1| 65951.65| 0|
| 15643966| null| null| null| null| null| 3| 143129.41| 2| 1| 64327.26| null|
| 15737452| null| null| Germany| Male| 58| null| 132602.88| 1| null| 5097.67| 1|
| 15788218| Henderson| 549| Spain|Female| null| 9| null| 2| 1| 14406.41| 0|
| 15661507| Muldrow| 587| Spain| Male| 45| 6| null| 1| 0| 158684.81| 0|
| 15568982| Hao| 726| France|Female| 24| 6| null| 2| 1| 54724.03| 0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```


4.4. Full join

04:58 PM (1s)

37

```
result3 = spark.sql("SELECT * FROM table1 t1 FULL OUTER JOIN table2 t2 ON t1.CustomerId = t2.CustomerId;")
result3.show()
```

(3) Spark Jobs

result3: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 11 more fields]

15574012	null	645	null	Male	44	15574012	8	113755.78	2	null	149756.71	1
15592389	null	684	France	Male	27	15592389	2	134603.88	1	1	71725.73	null
15592531	Bartlett	822	France	Male	null	15592531	7	null	2	1	10062.8	0
null	null	null	null	null	null	15600882	7	null	2	1	65951.65	0
null	null	null	null	null	null	15619304	8	159660.8	3	0	113931.57	1
null	null	null	null	null	null	15632264	10	null	2	null	26260.98	0
15634602	Hargrave	619	France	Female	42	15634602	2	null	1	1	101348.88	1
null	null	null	null	null	null	15643966	3	143129.41	2	1	64327.26	null
15647311	Hill	608	null	Female	41	15647311	1	83807.86	1	1	112542.58	0
15656148	Obinna	376	Germany	null	29	15656148	4	115046.74	4	0	119346.88	1
15661507	Muldrow	587	Spain	Male	45	15661507	6	null	1	0	158684.81	0
null	null	null	null	null	null	15691483	5	null	2	0	190857.79	null
15701354	null	null	France	null	null	15701354	1	null	2	0	93826.63	0
null	null	null	null	null	null	15737173	3	null	2	0	76390.01	0
15737452	null	null	Germany	Male	58	15737452	null	132602.88	1	null	5097.67	1
15737888	Mitchell	850	Spain	Female	null	15737888	2	125510.82	1	1	79084.1	null
null	null	null	null	null	null	15767821	null	102016.72	2	0	80181.12	0
15788218	Henderson	549	Spain	Female	null	15788218	9	null	2	1	14406.41	0
15792365	He	501	France	null	44	15792365	4	142051.07	2	1	74940.5	0

04:55 PM (1s)

38

```
full_join = df1.join(df2, on="CustomerId", how="outer")
full_join.show()
```

(3) Spark Jobs

full_join: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 10 more fields]

CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
15568982	Hao	726	France	Female	24	6	null	2	1	54724.03	0
15574012	null	645	null	Male	44	8	113755.78	2	null	149756.71	1
15592389	null	684	France	Male	27	2	134603.88	1	1	71725.73	null
15592531	Bartlett	822	France	Male	null	7	null	2	1	10062.8	0
15600882	null	null	null	null	null	7	null	2	1	65951.65	0
15619304	null	null	null	null	null	8	159660.8	3	0	113931.57	1
15632264	null	null	null	null	null	10	null	2	null	26260.98	0
15634602	Hargrave	619	France	Female	42	2	null	1	1	101348.88	1
15643966	null	null	null	null	null	3	143129.41	2	1	64327.26	null
15647311	Hill	608	null	Female	41	1	83807.86	1	1	112542.58	0
15656148	Obinna	376	Germany	null	29	4	115046.74	4	0	119346.88	1
15661507	Muldrow	587	Spain	Male	45	6	null	1	0	158684.81	0
15691483	null	null	null	null	null	5	null	2	0	190857.79	null
15701354	null	null	France	null	null	1	null	2	0	93826.63	0
15737173	null	null	null	null	null	3	null	2	0	76390.01	0
15737452	null	null	Germany	Male	58	null	132602.88	1	null	5097.67	1
15737888	Mitchell	850	Spain	Female	null	2	125510.82	1	1	79084.1	null
15767821	null	null	null	null	null	null	102016.72	2	0	80181.12	0

4.5. Left Semi

```

▶ 04:59 PM (1s) 40

result4 = spark.sql("SELECT * FROM table1 t1 WHERE EXISTS ( SELECT 1 FROM table2 t2 WHERE t1.CustomerId = t2.CustomerId );")
result4.show()

▶ (2) Spark Jobs

▶ result4: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 4 more fields]

+-----+-----+-----+-----+-----+-----+
|CustomerId| Surname|CreditScore|Geography|Gender| Age|
+-----+-----+-----+-----+-----+-----+
| 15634602| Hargrave|      619|   France|Female|  42|
| 15647311|   Hill|      608|    null|Female|  41|
| 15701354|   null|      null|   France|   null| null|
| 15737888| Mitchell|      850|   Spain|Female| null|
| 15574012|   null|      645|    null|Male|  44|
| 15592531| Bartlett|      822|   France|Male| null|
| 15656148| Obinna|      376| Germany|   null|  29|
| 15792365|   He|      501|   France|   null|  44|
| 15592389|   null|      684|   France|Male|  27|
| 15737452|   null|      null| Germany|Male|  58|
| 15788218|Henderson|      549|   Spain|Female| null|
| 15661507| Muldrow|      587|   Spain|Male|  45|
| 15568982|   Hao|      726|   France|Female|  24|
+-----+-----+-----+-----+-----+-----+

```

Left Semi join is not available in SQL but can be performed by using the sub queries concept.

```

▶ 04:55 PM (1s)

left_semi_join = df1.join(df2, on="CustomerId", how="leftsemi")
left_semi_join.show()

▶ (2) Spark Jobs

▶ left_semi_join: pyspark.sql.dataframe.DataFrame = [CustomerId: integer, Surname: string ... 4 more f

+-----+-----+-----+-----+-----+-----+
|CustomerId| Surname|CreditScore|Geography|Gender| Age|
+-----+-----+-----+-----+-----+-----+
| 15634602| Hargrave|      619|   France|Female|  42|
| 15647311|   Hill|      608|    null|Female|  41|
| 15701354|   null|      null|   France|   null| null|
| 15737888| Mitchell|      850|   Spain|Female| null|
| 15574012|   null|      645|    null|Male|  44|
| 15592531| Bartlett|      822|   France|Male| null|
| 15656148| Obinna|      376| Germany|   null|  29|
| 15792365|   He|      501|   France|   null|  44|
| 15592389|   null|      684|   France|Male|  27|
| 15737452|   null|      null| Germany|Male|  58|
| 15788218|Henderson|      549|   Spain|Female| null|
| 15661507| Muldrow|      587|   Spain|Male|  45|
| 15568982|   Hao|      726|   France|Female|  24|
+-----+-----+-----+-----+-----+-----+

```

5. Customers with Tenure less than 5 or CreditScore greater than 750:

```
05:17 PM (<1s) 43

result5 = spark.sql("SELECT * FROM credit WHERE Tenure < 5 OR CreditScore > 750;")
result5.show()

(1) Spark Jobs
result5: pyspark.sql.dataframe.DataFrame = [RowNumber: integer, CustomerId: integer ... 11 more fields]
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 4| 15701354| Boni| 699| France|Female| 39| 1| 0.0| 2| 0| 93826.63| 0|
| 5| 15737888| Mitchell| 850| Spain|Female| 43| 2|125510.82| 1| 1| 79084.1| 0|
| 7| 15592531| Bartlett| 822| France| Male| 50| 7| 0.0| 2| 1| 10062.8| 0|
| 8| 15656148| Obinna| 376| Germany|Female| 29| 4|115046.74| 4| 0| 119346.88| 1|
| 9| 15792365| He| 501| France| Male| 44| 4|142051.07| 2| 1| 74940.5| 0|
| 10| 15592389| H?| 684| France| Male| 27| 2|134603.88| 1| 1| 71725.73| 0|
| 12| 15737173| Andrews| 497| Spain| Male| 24| 3| 0.0| 2| 0| 76390.01| 0|
| 16| 15643966| Goforth| 616| Germany| Male| 45| 3|143129.41| 2| 1| 64327.26| 0|
| 17| 15737452| Romeo| 653| Germany| Male| 58| 1|132602.88| 1| 0| 5097.67| 1|
| 23| 15699309| Gerasimov| 510| Spain|Female| 38| 4| 0.0| 1| 0| 118913.53| 1|
| 24| 15725737| Mosman| 669| France| Male| 46| 3| 0.0| 2| 1| 8487.75| 0|
| 25| 15625047| Yen| 846| France|Female| 38| 5| 0.0| 1| 1| 187616.16| 0|
| 26| 15738191| Maclean| 577| France| Male| 25| 3| 0.0| 2| 1| 124508.29| 0|
| 27| 15736816| Young| 756| Germany| Male| 36| 2|136815.64| 1| 1| 170041.95| 0|
| 29| 15728693| McWilliams| 574| Germany|Female| 43| 3|141349.43| 1| 1| 100187.43| 0|
| 30| 15656300| Lucciano| 411| France| Male| 29| 0| 59697.17| 2| 1| 53483.21| 0|
| 31| 15589475| Azikiwe| 591| Spain|Female| 39| 3| 0.0| 3| 0| 140469.38| 1|
| 36| 15794171| Lombardo| 475| France|Female| 45| 0|134264.04| 1| 0| 27822.99| 1|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```
05:16 PM (1s) 44

filter2 = data_credit.filter((data_credit.Tenure < 5) | (data_credit.CreditScore > 750))
filter2.show()

(1) Spark Jobs
filter2: pyspark.sql.dataframe.DataFrame = [RowNumber: integer, CustomerId: integer ... 11 more fields]
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 4| 15701354| Boni| 699| France|Female| 39| 1| 0.0| 2| 0| 93826.63| 0|
| 5| 15737888| Mitchell| 850| Spain|Female| 43| 2|125510.82| 1| 1| 79084.1| 0|
| 7| 15592531| Bartlett| 822| France| Male| 50| 7| 0.0| 2| 1| 10062.8| 0|
| 8| 15656148| Obinna| 376| Germany|Female| 29| 4|115046.74| 4| 0| 119346.88| 1|
| 9| 15792365| He| 501| France| Male| 44| 4|142051.07| 2| 1| 74940.5| 0|
| 10| 15592389| H?| 684| France| Male| 27| 2|134603.88| 1| 1| 71725.73| 0|
| 12| 15737173| Andrews| 497| Spain| Male| 24| 3| 0.0| 2| 0| 76390.01| 0|
| 16| 15643966| Goforth| 616| Germany| Male| 45| 3|143129.41| 2| 1| 64327.26| 0|
| 17| 15737452| Romeo| 653| Germany| Male| 58| 1|132602.88| 1| 0| 5097.67| 1|
| 23| 15699309| Gerasimov| 510| Spain|Female| 38| 4| 0.0| 1| 0| 118913.53| 1|
| 24| 15725737| Mosman| 669| France| Male| 46| 3| 0.0| 2| 1| 8487.75| 0|
| 25| 15625047| Yen| 846| France|Female| 38| 5| 0.0| 1| 1| 187616.16| 0|
| 26| 15738191| Maclean| 577| France| Male| 25| 3| 0.0| 2| 1| 124508.29| 0|
| 27| 15736816| Young| 756| Germany| Male| 36| 2|136815.64| 1| 1| 170041.95| 0|
| 29| 15728693| McWilliams| 574| Germany|Female| 43| 3|141349.43| 1| 1| 100187.43| 0|
| 30| 15656300| Lucciano| 411| France| Male| 29| 0| 59697.17| 2| 1| 53483.21| 0|
| 31| 15589475| Azikiwe| 591| Spain|Female| 39| 3| 0.0| 3| 0| 140469.38| 1|
| 36| 15794171| Lombardo| 475| France|Female| 45| 0|134264.04| 1| 0| 27822.99| 1|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

6. Calculate the total number of active members and the average EstimatedSalary:

05:21 PM (<1s) 46

```
%sql
SELECT SUM(IsActiveMember) AS TotalActiveMembers, AVG(EstimatedSalary) AS AvgEstimatedSalary
FROM credit;
```

(2) Spark Jobs

_sqldf: pyspark.sql.dataframe.DataFrame = [TotalActiveMembers: long, AvgEstimatedSalary: double]

Table +

	1.2 TotalActiveMembers	1.2 AvgEstimatedSalary
1	5151	100090.2398809998

The above one done using SQL magic function in pyspark – Databricks platform.

05:20 PM (<1s) 47

```
agg1 = data_credit.agg({"IsActiveMember": "sum", "EstimatedSalary": "avg"}) \
.withColumnRenamed("sum(IsActiveMember)", "TotalActiveMembers") \
.withColumnRenamed("avg(EstimatedSalary)", "AvgEstimatedSalary")
agg1.show()
```

(2) Spark Jobs

agg1: pyspark.sql.dataframe.DataFrame = [TotalActiveMembers: long, AvgEstimatedSalary: double]

```
+-----+-----+
|TotalActiveMembers|AvgEstimatedSalary|
+-----+-----+
|          5151| 100090.2398809998|
+-----+-----+
```

Thank you

Sivaprakash V